16-720 Computer Vision: Homework 2
Keypoints - Detectors, Descriptors and Matching

Instructor: Srinivasa Narasimhan
TAs: Minh Vo, Supreeth Achar

Due: Monday, February 24, 11:59:59 p.m.

Introduction

In this homework, you will implement an interest point detector and feature descriptor. Interest point detectors find particularly salient points in an image upon which we can extract a feature descriptor. SIFT, SURF and BRIEF are all examples of descriptors. In our case, we will be using BRIEF. Once we have extracted the interest points, we can use descriptors to match them between images to do neat things like panorama stitching or scene reconstruction.

BRIEF is one the simplest feature descriptors to implement. It has a very compact representation, is quick to compute, and has a discriminative yet easily computed distance metric. This allows for real-time computation, as you have seen in class. Most importantly, as you will see, it is also just as powerful as more complex descriptors like SIFT, for many cases.

1 Keypoint Detector

For our implementation, we will be using the detector similar to the one introduced in class. A good reference for its implementation can be found in [3]. Keypoints are found by using the Difference of Gaussian (DoG) detector. This detector finds points that are extrema in both scale and space of a DoG pyramid. This is described in [1], an important paper in our field.

Here, we will be implementing a simplified version of the DoG detector described in Section 3 of [3]. The parameters you will use for the following sections are $\sigma_0 = 1$, $k = \sqrt{2}$, $levels = [-1\ 0\ 1\ 2\ 3\ 4]$, $th_{contrast} = 0.03$ and $th_r = 12$.

1.1 Gaussian Pyramid

In order to create a DoG pyramid, we will first need to create a Gaussian pyramid. Gaussian pyramids are constructed by progressively applying a lowpass filter to the input image. We provide you this function.

```
GaussianPyramid = createGaussianPyramid(im, sigma0, k, levels)
```

The function takes as input a grayscale image im with values between 0 and 1 (hint: `im2double`), the scale of the zeroth level of the pyramid `sigma0`, the pyramid factor `k`, and a vector levels specifying the levels of the pyramid to construct.

At level l in the pyramid, the image is smoothed by a Gaussian filter with $\sigma_l = \sigma_0 k^l$. `GaussianPyramid` is a $R \times C \times L$ matrix, where $R \times C$ is the size of the input image `im` and and L is the size of levels. An example of a Gaussian pyramid can be seen in Figure 1.1.

Figure 1: Example Gaussian pyramid for model `chickenbroth.jpg`

Figure 2: Example DoG pyramid for model `chickenbroth.jpg`. Images scaled for display.

1.1.1 The DoG Pyramid (5 pts)

The DoG pyramid is obtained by subtracting successive levels of the Gaussian pyramid.

$$D_l(x, y, \sigma_l) = (G(x, y, \sigma_{l-1} - G(x, y, \sigma_l)) * I(x, y) \tag{1}$$

Write the following function to construct a Difference of Gaussian pyramid:

```
[DoGPyramid, DoGLevels] = createDoGPyramid(GaussianPyramid, levels)
```

The function should return `DoGPyramid` an $R \times C \times (L-1)$ matrix of the DoG pyramid created using `GaussianPyramid`. Note that you will have one level less than the Gaussian Pyramid. `DoGLevels` is an $(L-1)$ vector specifying the corresponding levels of the DoG Pyramid. An example of the DoG pyramid can be seen in Figure 1.1.1.

1.1.2 Edge Suppression (10 pts)

The Difference of Gaussian function responds strongly on corners and edges in addition to blob-like objects. However, edges are not desirable for feature extraction as they are not distinctive and do not provide stable localization of the keypoint. Here, we will implement the edge removal method described in Section 4.1 of [3], which is based on the principal curvature ratio in a local neighborhood of a point. The paper makes the observation that edge points will have a "large principal curvature across the edge but a small one in the perpendicular direction."

Implement the following function:

```
PrincipalCurvature = computePrincipalCurvature(DoGPyramid)
```

2

Figure 3: Left, keypoints with edge suppression, right, without for model `chickenbroth.jpg`

The function takes in `DoGPyramid` generated in the previous section and returns `PrincipalCurvature`, a matrix of the same size where each point contains the curvature ratio R for the corresponding point in the DoG pyramid:

$$R = \frac{Tr(H)^2}{Det(H)} = \frac{(\lambda_{min} + \lambda_{max})^2}{\lambda_{min}\lambda_{max}} \tag{2}$$

Here, H is the Hessian of the Difference of Gaussian function (i.e. one level of the DoG pyramid) computed by using pixel differences as mentioned in Section 4.1 of [3]. (hint: Matlab function `gradient`). This is similar in spirit but different from the Harris matrix you saw in class. Both methods examine the eigen values of a matrix, but the method in [3] performs a test that does not require direct computation of the values.

We can see that R reaches its minimum when the two eigenvalues λ_{min} and λ_{max} are equal, meaning that the curvature is the same in the two principal directions. Edge points, in general, will have a principal curvature significantly larger in one direction than the other. Thus, to remove edge points, we can check that $R < thr$ for some thr. Figure 1.1.2 shows the DoG detector with and without edge suppression.

1.1.3 Detecting Extrema (10 pts)

To detect corner-like keypoints that are scale-invariant, the DoG detector chooses points that are local extrema in both scale and space. Here, we will consider only a point's eight neighbors in space and its two neighbors in scale (one in the scale above and one in the scale below). Write the function:

```
locs = getLocalExtrema(DoGPyramid, DoG levels, PrincipalCurvature,
                       th_contrast, th_r)
```

3

This function takes as input `DoGPyramid` and `DoG_levels` from Section 1.1.1 and PrincipalCurvature from Section 1.1.2. It also takes two threshold values, `th_contrast` and `th_r`. The threshold th contrast should remove any point that is a local extremum but does not have a Difference of Gaussian response magnitude above this threshold (i.e. $|D(x, y, \sigma| > th_{contrast})$. The threshold `th_r` should remove any edge-like points that have too large a principal curvature ratio specified by `PrincipalCurvature`.

The function should return `locs`, an $N \times 3$ matrix where the DoG pyramid achieves a local extrema in both scale and space, and also satisfies the two thresholds. The first and second column of locs should be the (x, y) values of the local extremum and the third column should contain the corresponding level of the DoG pyramid where it was detected. (Try to eliminate loops in the function so that it runs efficiently.)

1.2 Putting it together (5 pts)

Write the following function to combine the above parts into a DoG detector:

```
[locs, GaussianPyramid] = DoGdetector(im, sigma0, k, levels,th_contrast,
                            th_r)
```

The function should take in a gray scale image, `im`, scaled between 0 and 1, and the parameters `sigma0`, `k`, `levels`, `th_contrast`, and `th_r`. It should call each of the above functions and return the keypoints in locs and the Gaussian pyramid in `GaussianPyramid`. Figure 1.1.2 left shows the keypoints detected for an example image. Note that we are dealing with real images here, so your keypoint detector may find points with high scores that you do not perceive to be corners.

2 BRIEF Descriptor

Now that we have keypoints that tell us where to find the most informative points in the image, we can compute descriptors that can be used to match to other views of the same point in different images. The BRIEF descriptor encodes information from a 9×9 patch p centered around the interest point at the *characteristic scale* of the interest point. See the lecture notes for Point Feature Detectors if you need to refresh your memory.

2.1 Creating a Set of BRIEF Tests (5 pts)

The descriptor itself is a vector that is n-bits long, where each bit is the result of the following simple test:

$$\tau(p; x, y) := \begin{cases} 1, & \text{if } p(x) < p(y). \\ 0, & \text{otherwise.} \end{cases} \tag{3}$$

Set n to 256 bits. There is no need to encode the test results as actual bits. It is fine to encode them as a 256 element vector.

There are many choices for the 256 test pairs. The authors describe and test some of them in [2]. Read section 3.2 of that paper and implement one of these solutions. You should generate a static set of test pairs and save that data to a file. You will use these pairs for all subsequent computations of the BRIEF descriptor. Please include this file in your submission so that we do not need to regenerate your test pattern.

Write the function:

```
[compareX, compareY] = makeTestPattern(patchWidth, nbits)
```

Where `patchWidth` is the width of the image patch (usually 9) and `nbits` is the number of tests in the BRIEF descriptor. `compareX` and `compareY` are linear indices into the $patchWidth \times patchWidth$ image patch and are each $nbits \times 1$ vectors. Run this routine for the given parameters and save the results in `testPattern.mat`.

2.2   Compute the BRIEF Descriptor (10 pts)

It is now time to compute the BRIEF descriptor for the detected keypoints. Write the function:

```
[locs,desc] = computeBrief(im, locs, levels, compareX, compareY)
```

Where im is grayscale image with values from 0 to 1, locs are the keypoint locations returned by the DoG detector from Section 1.2, levels are the scale levels that were given in Section 1, and `compareX` and `compareY` are the test patterns computed in Section 2.1.

The function returns `locs`, an $m \times 3$ vector, where the first two columns are the image coordinates of keypoints and the third column is the pyramid level of the keypoints, and `desc`, an $m \times nbits$ matrix of stacked BRIEF descriptors. m is the number of valid descriptors in the image and will vary.

2.3   Putting it all Together (5 pts)

Write a function

```
[locs, desc] = brief(im)
```

Which accepts a grayscale image im with values between zero and one and returns `locs`, an $m \times 3$ vector, where the first two columns are the image coordinates of keypoints and the third column is the pyramid level of the keypoints, and `desc`, an $m \times nbits$ matrix of stacked BRIEF descriptors. m is the number of valid descriptors in the image and will vary.

This function should perform all the necessary steps to extract the descriptors from the image, including

- Load parameters and test patterns

- Get keypoint locations

- Compute a set of valid BRIEF descriptors

2.4   Descriptor Matching (15 pts)

A descriptor's strength is in its ability to match to other descriptors generated by the same world point, despite change of view, lighting, etc. The distance metric used to compute the similarity between two descriptors is critical. For BRIEF, this distance metric is the Hamming distance. The Hamming distance is simply the number of bits in two descriptors that differ. (Note that the position of the bits matters.) Write the function

```
[matches] = briefMatch(desc1, desc2, ratio)
```

Figure 4: Example of BRIEF matches for model `chickenbroth.jpg` and `chickenbroth 01.jpg`.

Which accepts an $m \times nbits$ stack of BRIEF descriptors from a first image and a $n \times nbits$ stack of BRIEF descriptors from a second image and returns a $p \times 2$ matrix of matches, where the first column are indices into `desc1` and the second column are indices into `desc2`. Note that m, n, and p may be different sizes.

Section 7.2 of [3] introduces a ratio test to suppress matches between descriptors that are not particularly discriminative. Implement this test and select a value for ratio. In your PDF, mention how you designed this test and how you selected ratio.

Write a test script `testMatch` to load two of the chickenbroth images, compute feature matches, and use the provided `plotMatches` function to visualize the result.

```
plotMatches(im1, im2, matches, locs1, locs2)
```

Where `im1` and `im2` are grayscale images from 0 to 1, matches is the list of matches returned by `briefMatch` and `locs1` and `locs2` are the locations of keypoints from brief.

Save the resulting figure and submit it in your PDF. Figure 2.4 is an example result. A good test is to check that you can match an image to itself.

2.5   BRIEF and rotations (10 pts)

Take one of the test images and match it to itself while rotating the second image in increments of 10 degrees. Count the number of correct matches at each rotation and construct a plot. Include this in your PDF and explain why you think the descriptor behaves this way. Create a script `briefRotTest` that performs this task.

3   What to Submit

Your submission should consist of only 1 file, a zip file named `<AndrewId>.zip`. This zip file should contain

- a folder `matlab` containing all the `.m` and `.mat` files you were asked to write

- a pdf named `writeup.pdf` containing the results, explanations and images asked for in the assignment

You may leave the `data` folder and the provided `.m` files in your submission, but they are not needed. Please submit you homework using blackboard as usual.

References

[1] P. Burt and E. Adelson. The Laplacian Pyramid as a Compact Image Code. *IEEE Transactions on Communications*, 31(4):532–540, April 1983.

[2] Michael Calonder, Vincent Lepetit, Christoph Strecha, and Pascal Fua. BRIEF : Binary Robust Independent Elementary Features .

[3] David G. Lowe. Distinctive Image Features from Scale-Invariant Keypoints. *International Journal of Computer Vision*, 60(2):91–110, November 2004.